

FROM DATABASES TO ARCHITECTURES FOR BIG DATA MANAGEMENT

DATABASE FUNDAMENTALS (RECALL/CRASH COURSE)

DATABASE FUNDAMENTALS (RECALL/INTRODUCTION)

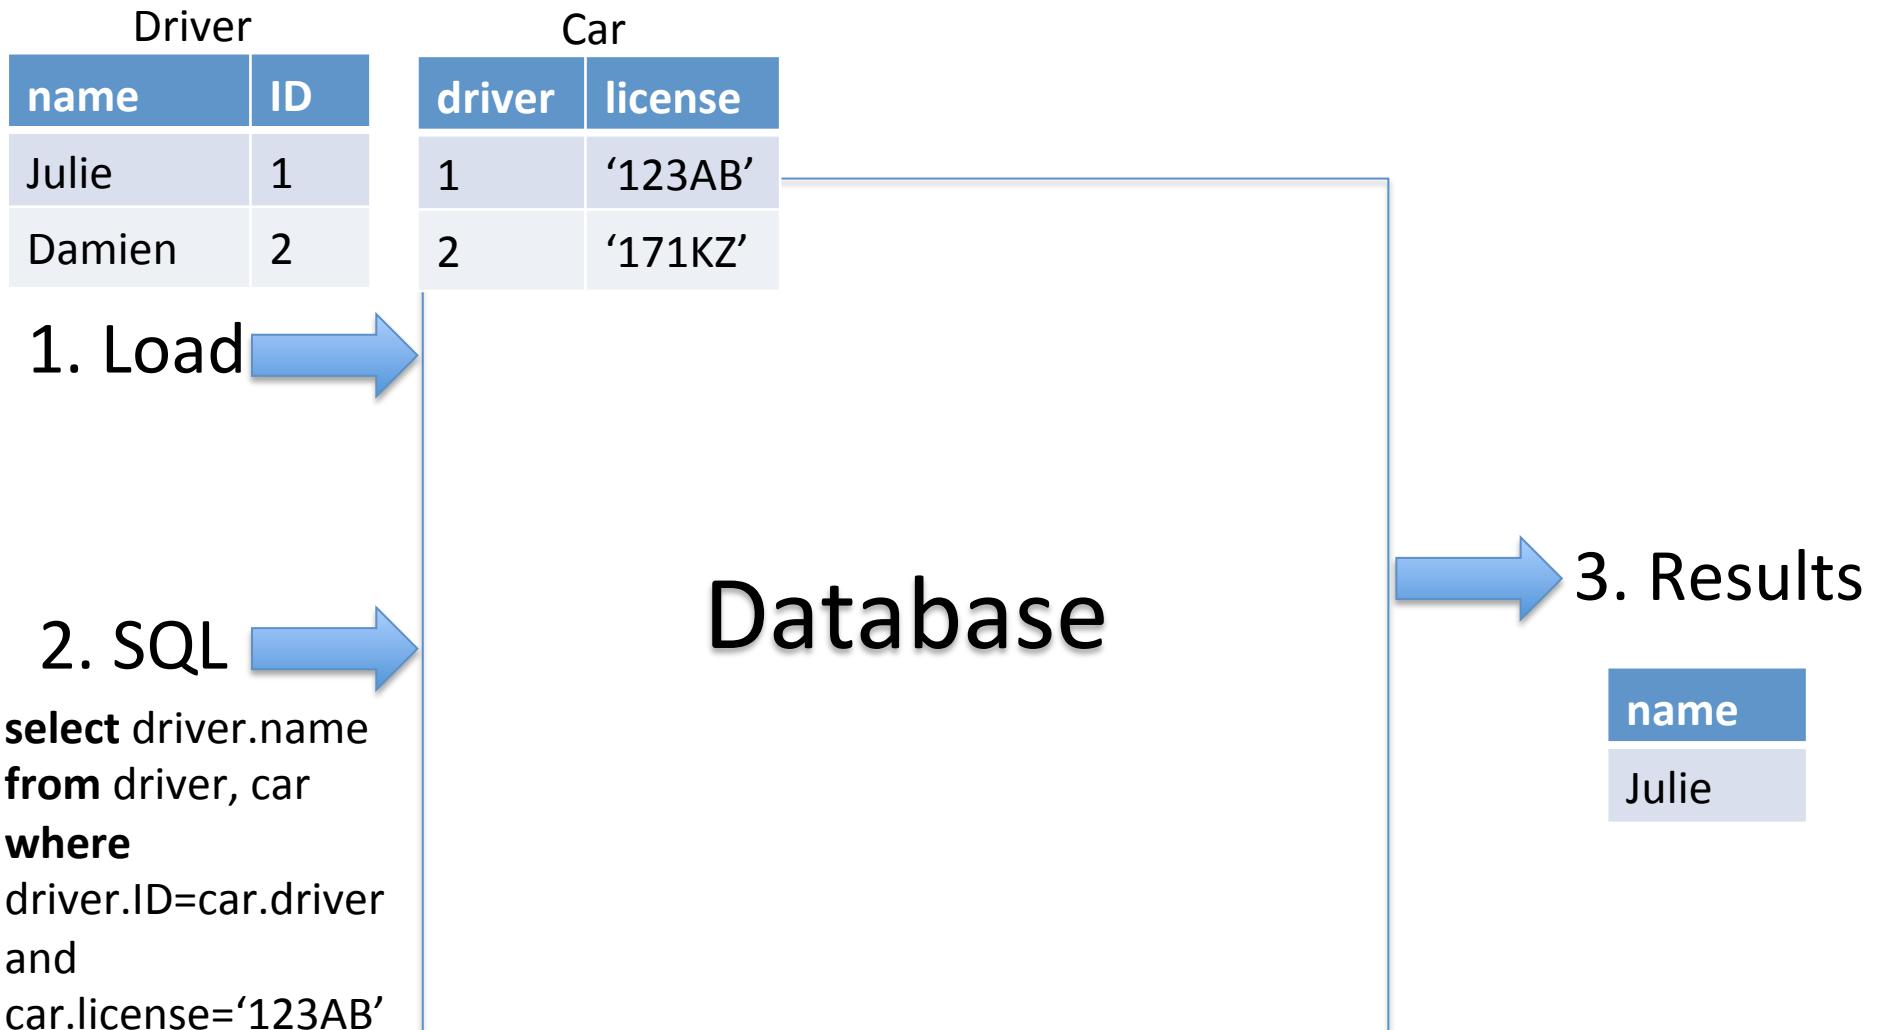
Database internals

(We illustrate for relational databases, as they are the most mature)

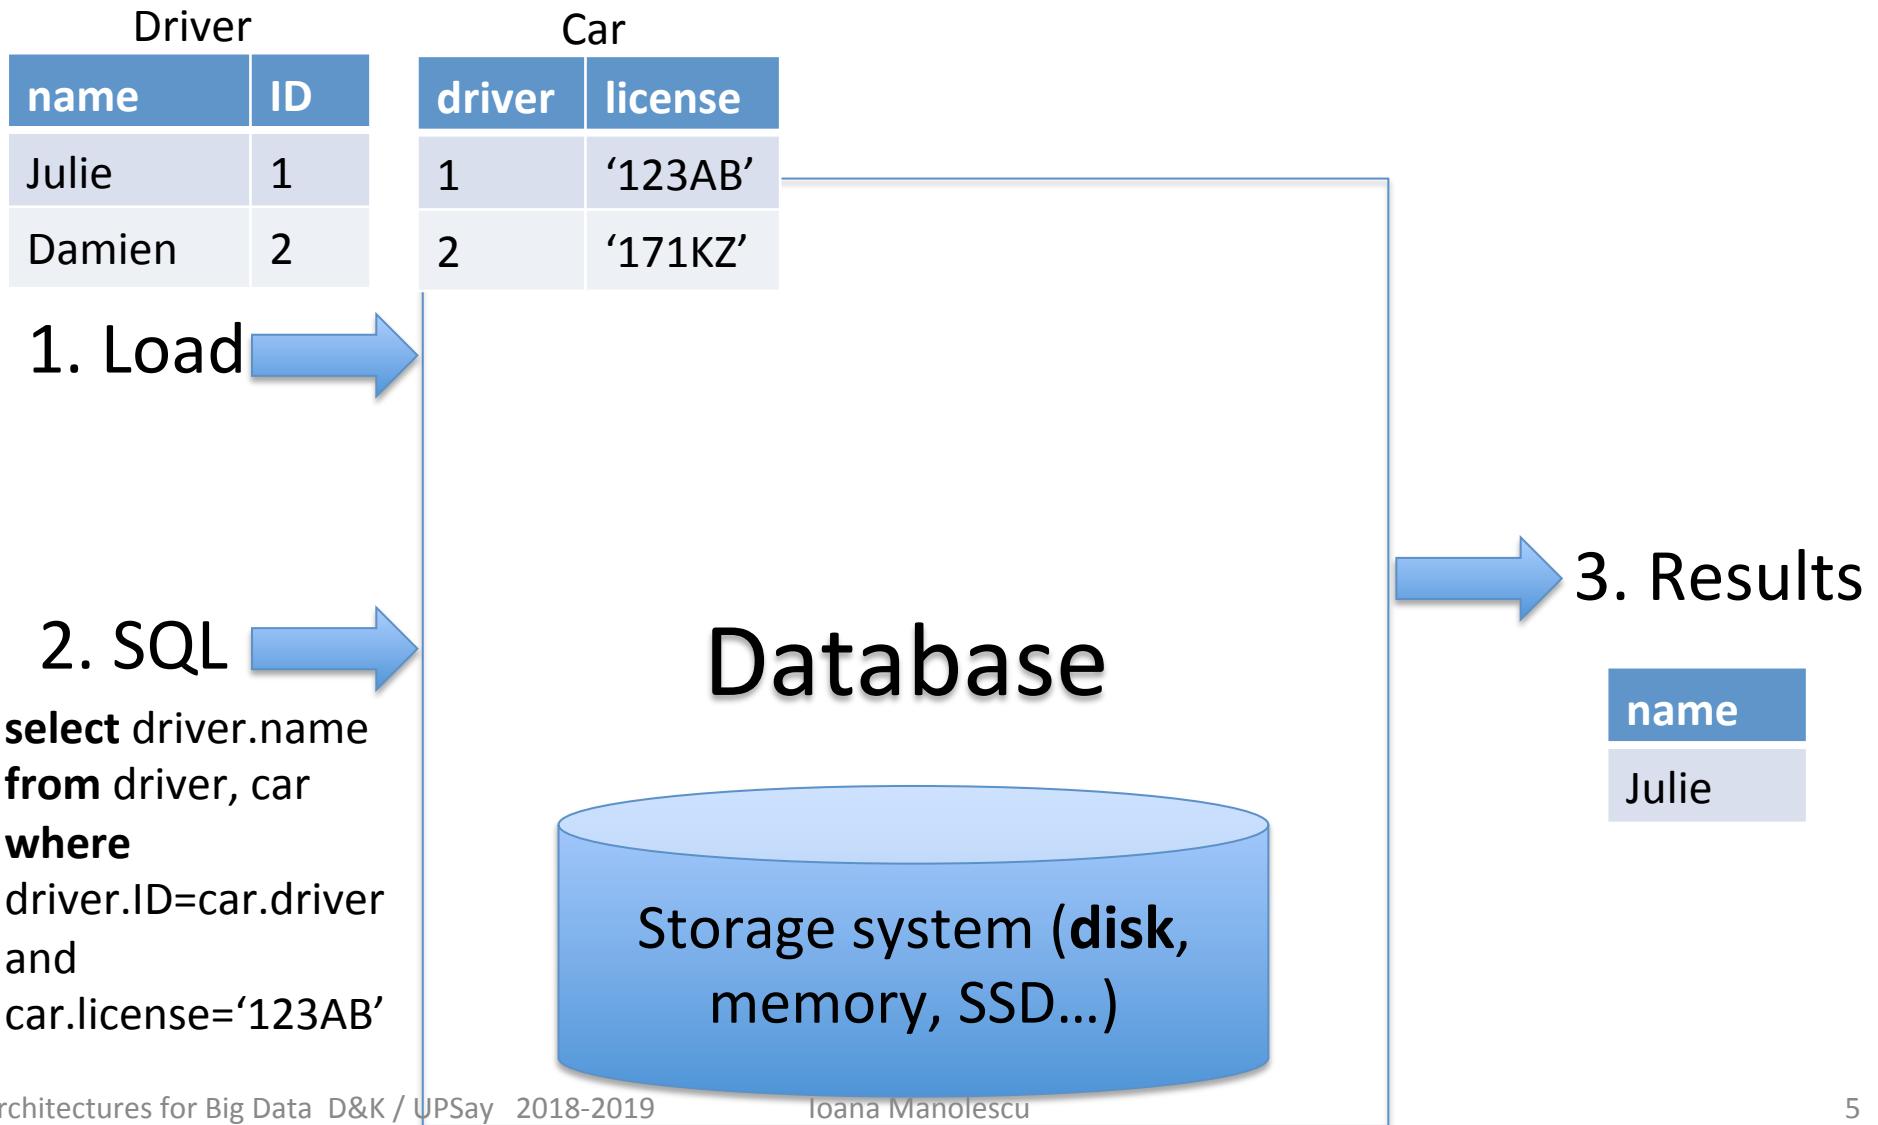
What's in a database?



What's in a database?



What's in a database?



What's in a database?

1. Load

2. SQL

```
select driver.name  
from driver, car  
where  
driver.ID=car.driver  
and  
car.license='123AB'
```

Database

Driver		Car	
name	ID	driver	license
Julie	1	1	'123AB'
Damien	2	2	'171KZ'

3. Results

name

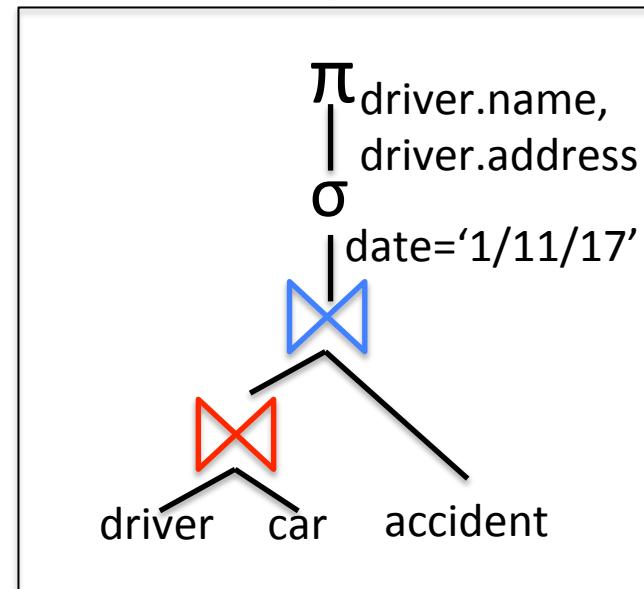
Julie

What's in a database?

SQL

```
select driver.name,  
       driver.address  
  from driver, car,  
       accident  
 where  
       driver.ID=car.driver  
       and  
       car.license=accident  
       .carLicense and  
       accident.date='1/11  
       /17'
```

select... from driver, car, accident where...



Driver		Accident		Car	
name	ID	driver	date	license	ID
Julie	1	1	'1/11/17'	'123AB'	
Damien	2	2		'171KZ'	

Query language

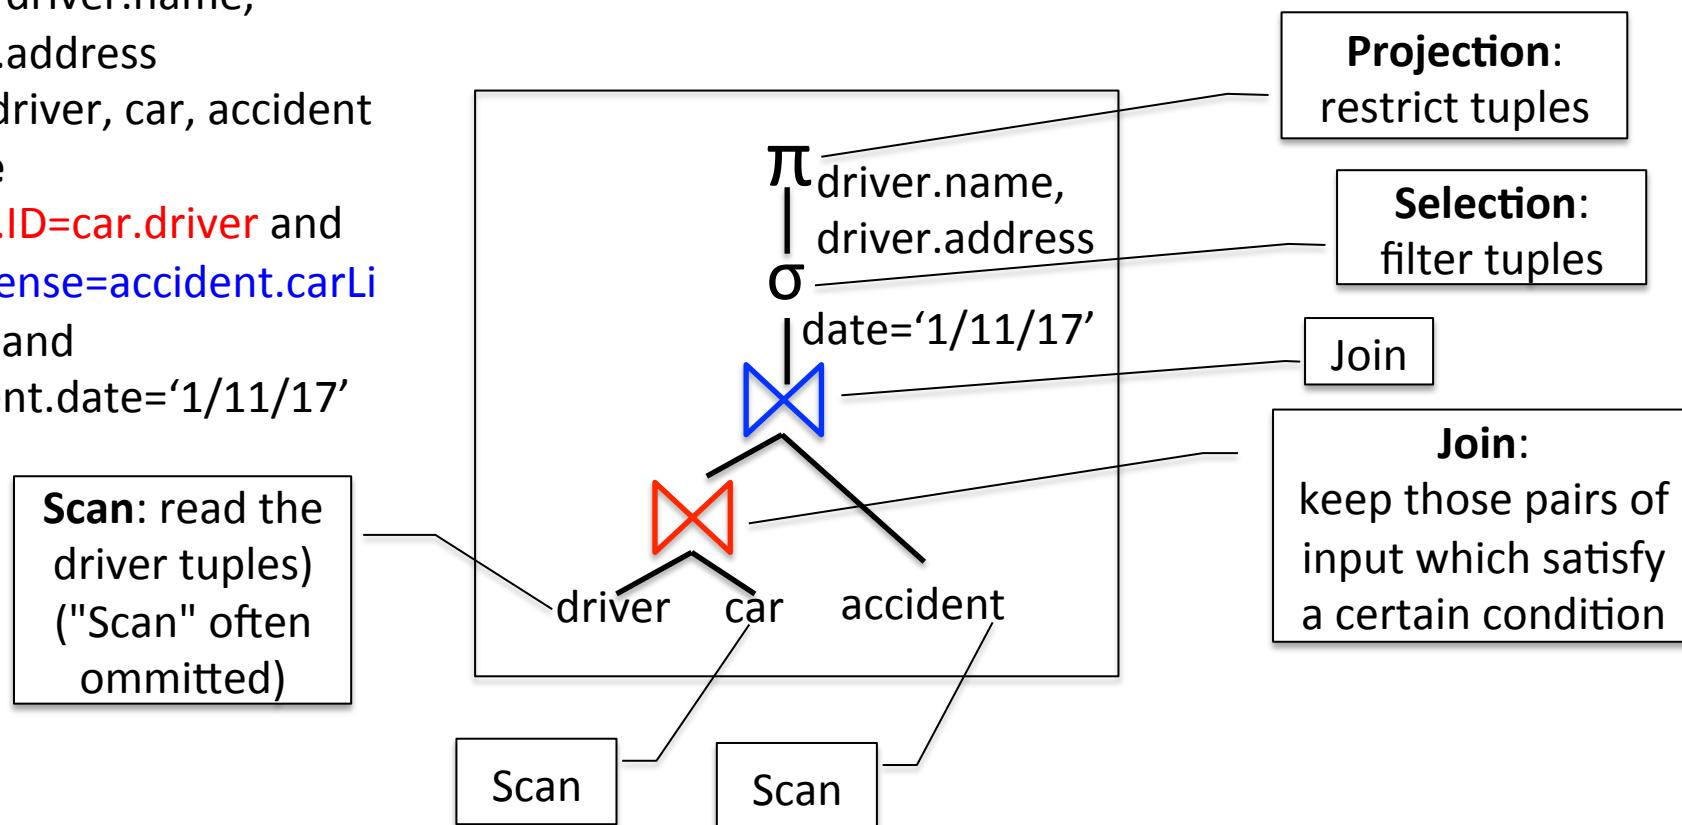
Logical plan

Logical query plans

- Trees made of logical operators, each of which specializes in a certain task

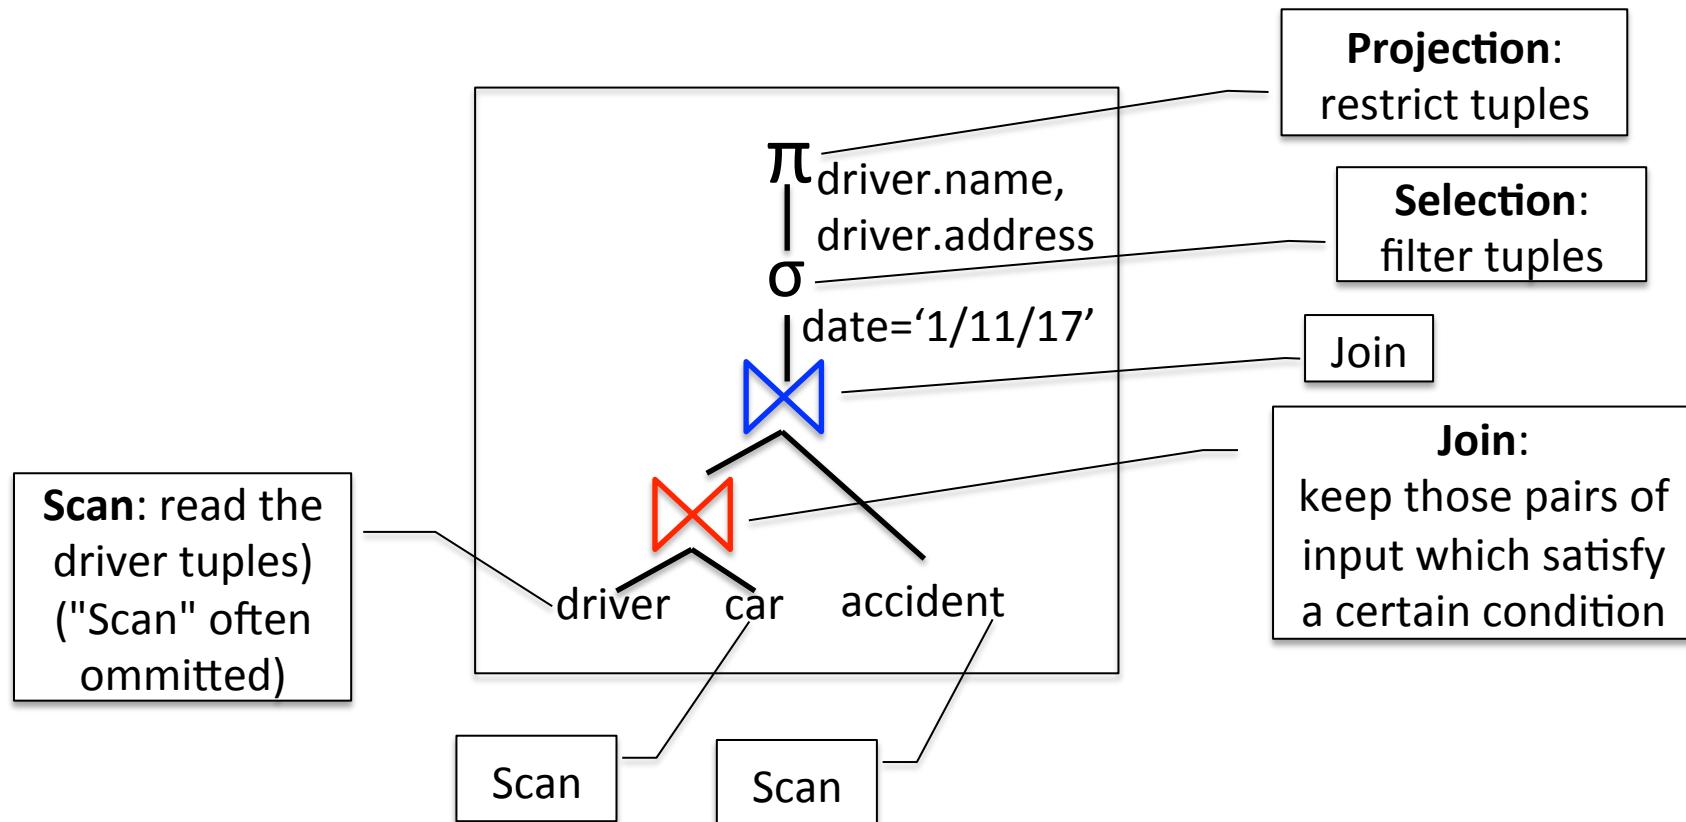
SQL:

```
select driver.name,  
       driver.address  
  from driver, car, accident  
 where  
   driver.ID=car.driver and  
   car.license=accident.carLi  
   cense and  
   accident.date='1/11/17'
```



Logical query plans

- Trees made of logical operators, each of which specializes in a certain task
- Logical operators: they are defined by their result, not by an algorithm
- Physical operators (see next) implement actual algorithms



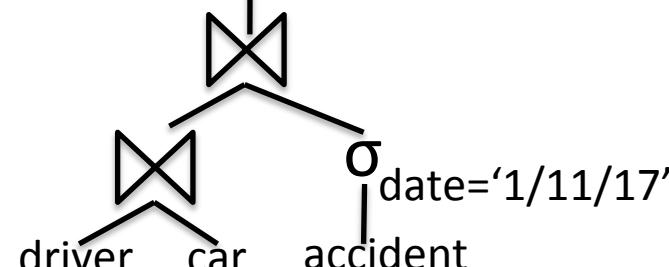
What's in a database?

SQL

```
select driver.name,  
       driver.address  
  from driver, car,  
       accident  
 where  
       driver.ID=car.driver  
       and  
       car.license=accident  
       .carLicense and  
       accident.date='1/11  
       /13'
```

select... from driver, car, accident where...

$\pi_{\text{driver.name},\text{driver.address}}$



A diagram of a relational table with three columns: `Driver`, `Accident`, and `Car`. The `Driver` column contains the attributes `name` and `ID`. The `Accident` column contains the attribute `driver`. The `Car` column contains the attribute `license`. The table has two rows. The first row contains `Julie` and `1` in the `Driver` column, and `1` in the `Accident` column. The second row contains `Damien` and `2` in the `Driver` column, and `2` in the `Accident` column. The `Car` column is empty for both rows.

Driver		Accident	Car
<code>name</code>	<code>ID</code>	<code>driver</code>	<code>license</code>
Julie	1	1	'123AB'
Damien	2	2	'171KZ'

Query language

Logical plan 1

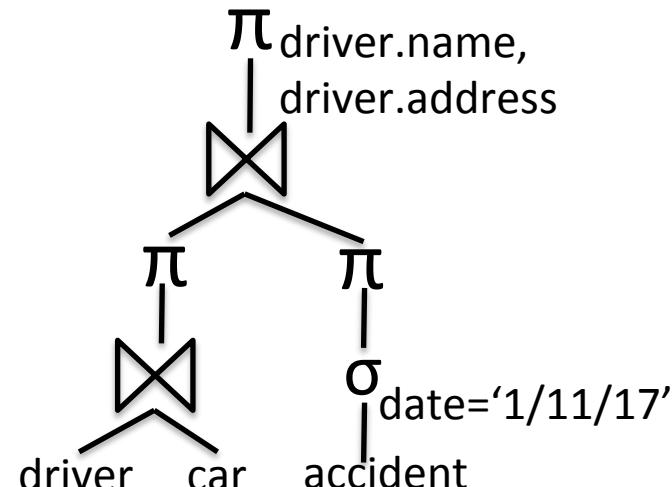
Logical plan 2

What's in a database?

SQL

```
select driver.name,  
       driver.address  
  from driver, car,  
       accident  
 where  
       driver.ID=car.driver  
       and  
       car.license=accident  
       .carLicense and  
       accident.date='1/11  
       /17'
```

select... from driver, car, accident where...



Driver		Accident		Car	
name	ID	driver	date	license	
Julie	1	1	'1/11/17'	'123AB'	
Damien	2	2		'171KZ'	

Query language

Logical plan 1

Logical plan 2

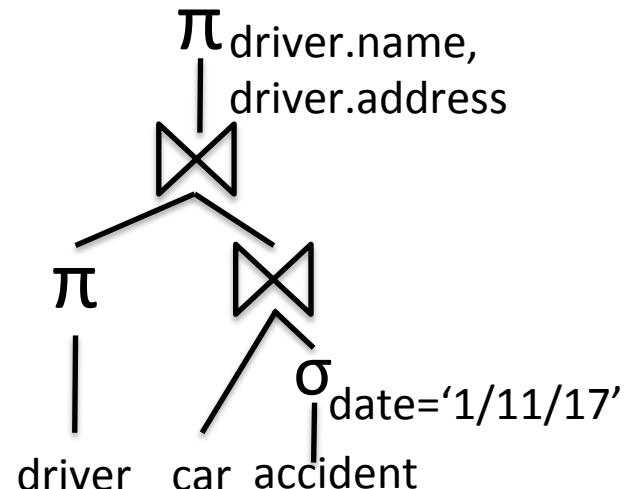
Logical plan 3

What's in a database?

SQL

```
select driver.name,  
       driver.address  
  from driver, car,  
       accident  
 where  
       driver.ID=car.driver  
       and  
       car.license=accident  
       .carLicense and  
       accident.date='1/11  
       /17'
```

select... from driver, car, accident where...



A diagram of a relational database table. It has three columns: `Driver`, `Accident`, and `Car`. The `Driver` column contains attributes `name` and `ID`. The `Accident` column contains attribute `driver`. The `Car` column contains attribute `license`. The table has two rows. The first row contains `Julie` in the `name` column, `1` in the `ID` column, and `1` in the `driver` column. The second row contains `Damien` in the `name` column, `2` in the `ID` column, and `2` in the `driver` column. The `license` column is empty for both rows.

Driver		Accident	Car
<code>name</code>	<code>ID</code>	<code>driver</code>	<code>license</code>
Julie	1	1	'123AB'
Damien	2	2	'171KZ'

Query language

Logical plan 1

Logical plan 2

Logical plan 3

Logical plan 4

Results

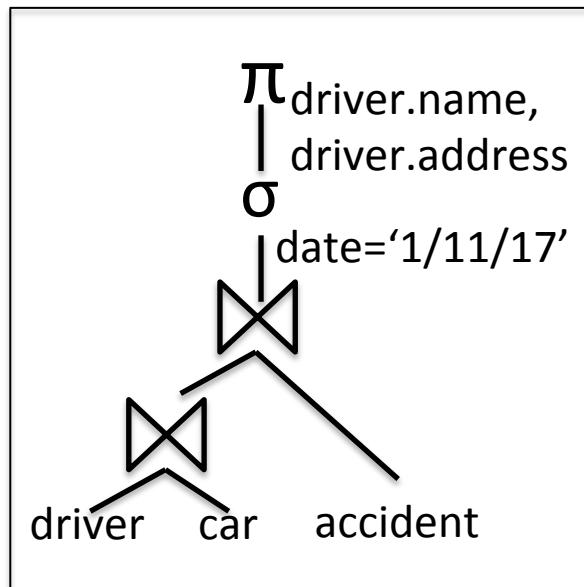
Logical query optimization

- Enumerates logical plans
- All logical plans compute the query result
 - They are **equivalent**
- Some are (much) more **efficient** than others
- **Logical optimization**: moving from a plan to a more efficient one
 - Pushing selections
 - Pushing projections
 - Join reordering: most important source of optimizations

Logical query optimization example

1.000.000 cars, 1.000.000 drivers, 1.000 accidents, 2 cars per accident, 10 accidents on 1/11/17

« Name and address of drivers in accidents on 1/11/2017? »



Cost of an operator: depends on the number of tuples (or tuple pairs) which it must process
e.g. $c_{\text{disk}} \times$ number of tuples read from disk
e.g. $c_{\text{cpu}} \times$ number of tuples compared

Cardinality of an operator's output: how many tuples result from this operator

The cardinality of one operator's output determines the cost of its parent operator

Plan **cost** = the sum of the costs of all operators in a plan

Logical query optimization example

1.000.000 cars, 1.000.000 drivers, 1.000 accidents, 2 cars per accident, 10 accidents on 1/11/17

« Name and address of drivers in accidents on 1/11/2017? »

Scan **costs**: $cs \times (10^6 + 10^6 + 10^3)$

Scan cardinality estimations: 10^6 , 10^6 , 10^3

Driver-car join **cost** estimation: $cj \times (10^6 \times \overline{10^6} = 10^{12})$

Driver-car join cardinality estimation: 10^9

Driver-car-accident join **cost** estim.: $ci \times (10^9 \times 10^3 = 10^{12})$

Driver-car-accident join cardinality estimation: 2×10^3

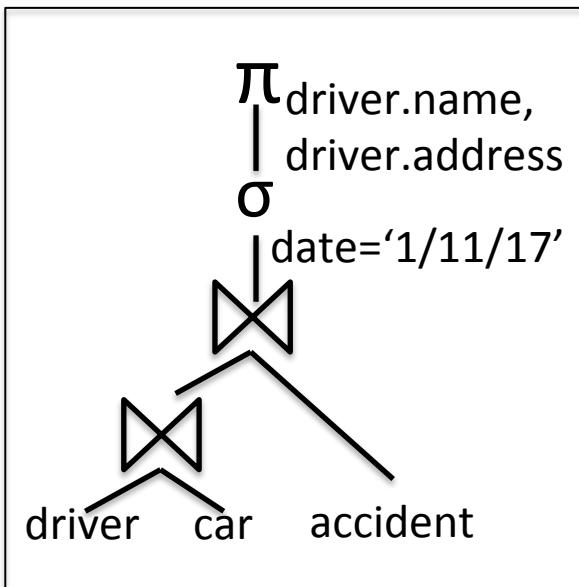
Selection **cost** estimation: $cf \times (2 \times 10^3)$

Selection cardinality estimation: 10

Projection (similar), negligible

Total **cost** estimation: $10 + 3 \times 10^3 + 2 \times 10^6 + 10^{12} \sim 10^{12}$

Pessimistic (worst-case) estim.



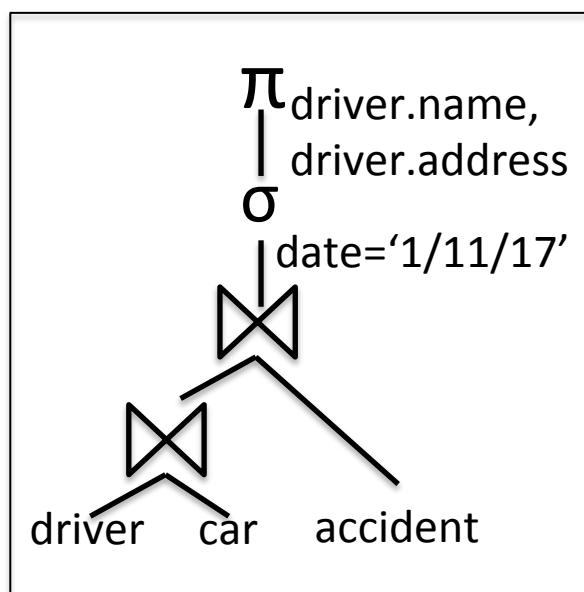
cs, cj, cf constant

Logical query optimization example

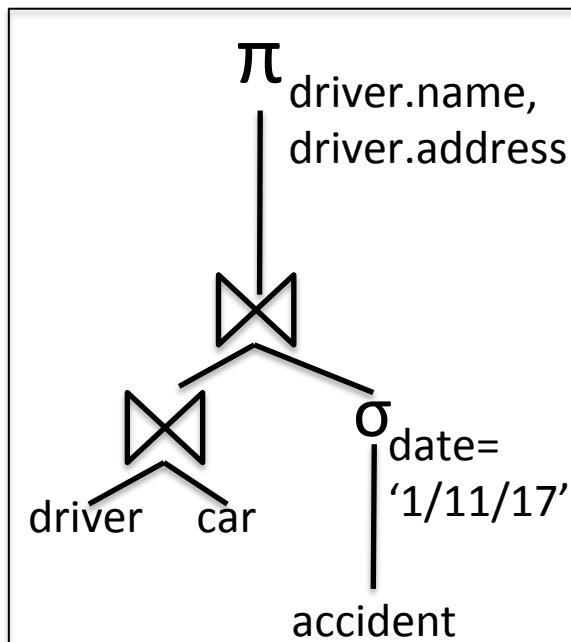
1.000.000 cars, 1.000.000 drivers, 1.000 accidents, 2 cars per accident, 10 accidents on 1/11/17

« Name and address of drivers in accidents on 1/11/2017? »

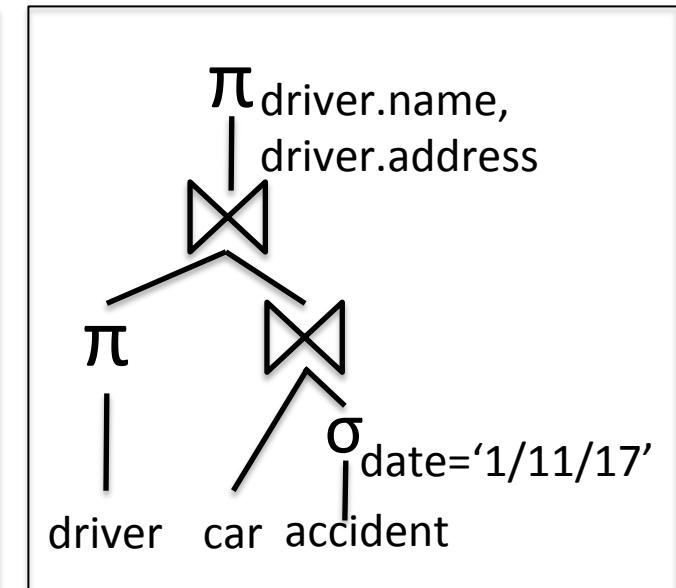
Three plans, same scan costs (neglected below); join costs dominant



$10^9 + 10^{12} \sim 10^{12}$



$10^9 + 10^7 \sim 10^9$



$10^7 + 2*10^7 \sim 3*10^7$

Logical query optimization example

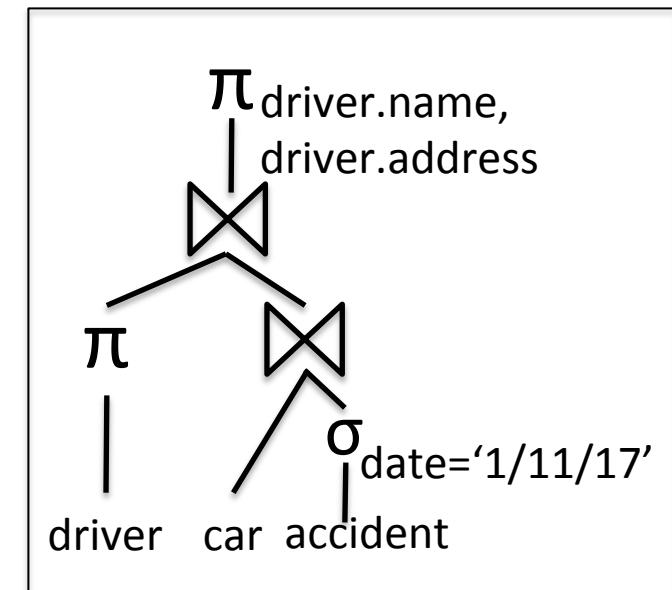
1.000.000 cars, 1.000.000 drivers, 1.000 accidents, 2 cars per accident, 10 accidents on 1/11/17

« Name and address of drivers in accidents on 1/11/2017? »

Three plans, same scan costs (neglected below); join costs dominant

The best plan reads only the accidents that have to be consulted

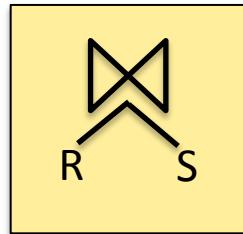
- **Selective data access**
- Typically supported by an **index**
 - Auxiliary data structure, built on top of the data collection
 - Allows to access directly objects satisfying a certain condition



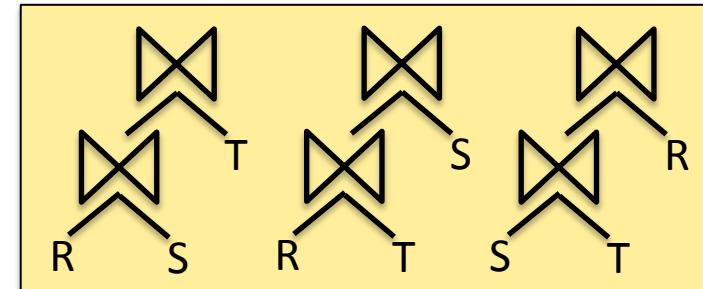
$$10^7 + 2*10^7 \sim 3*10^7$$

Join ordering is the main problem in logical query optimization

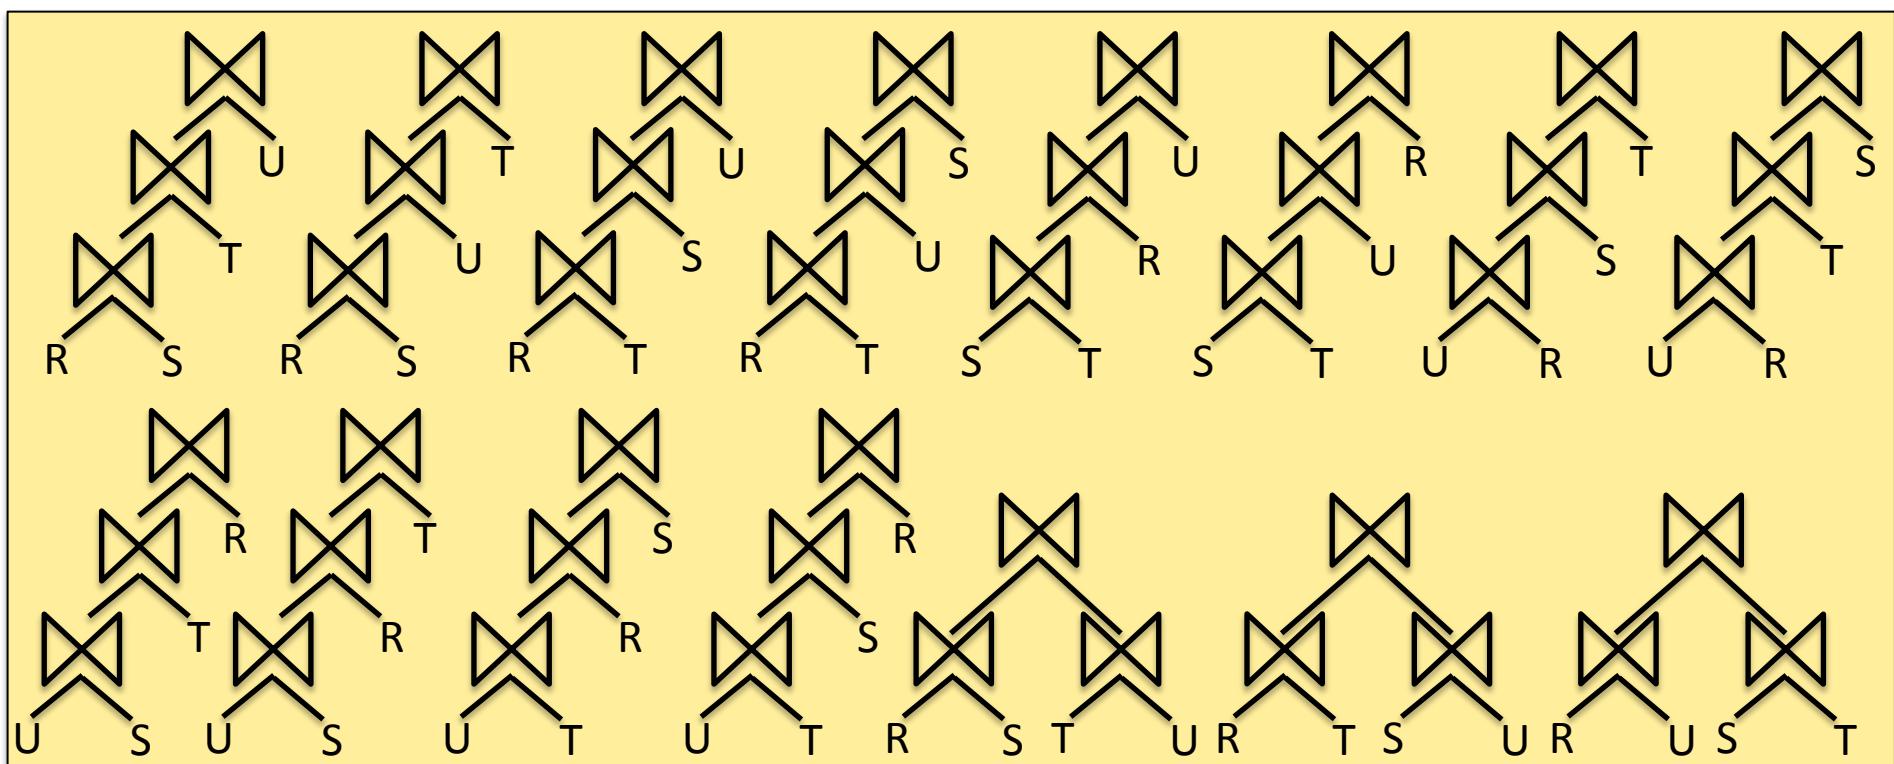
N=2:



N=3:



N=4:



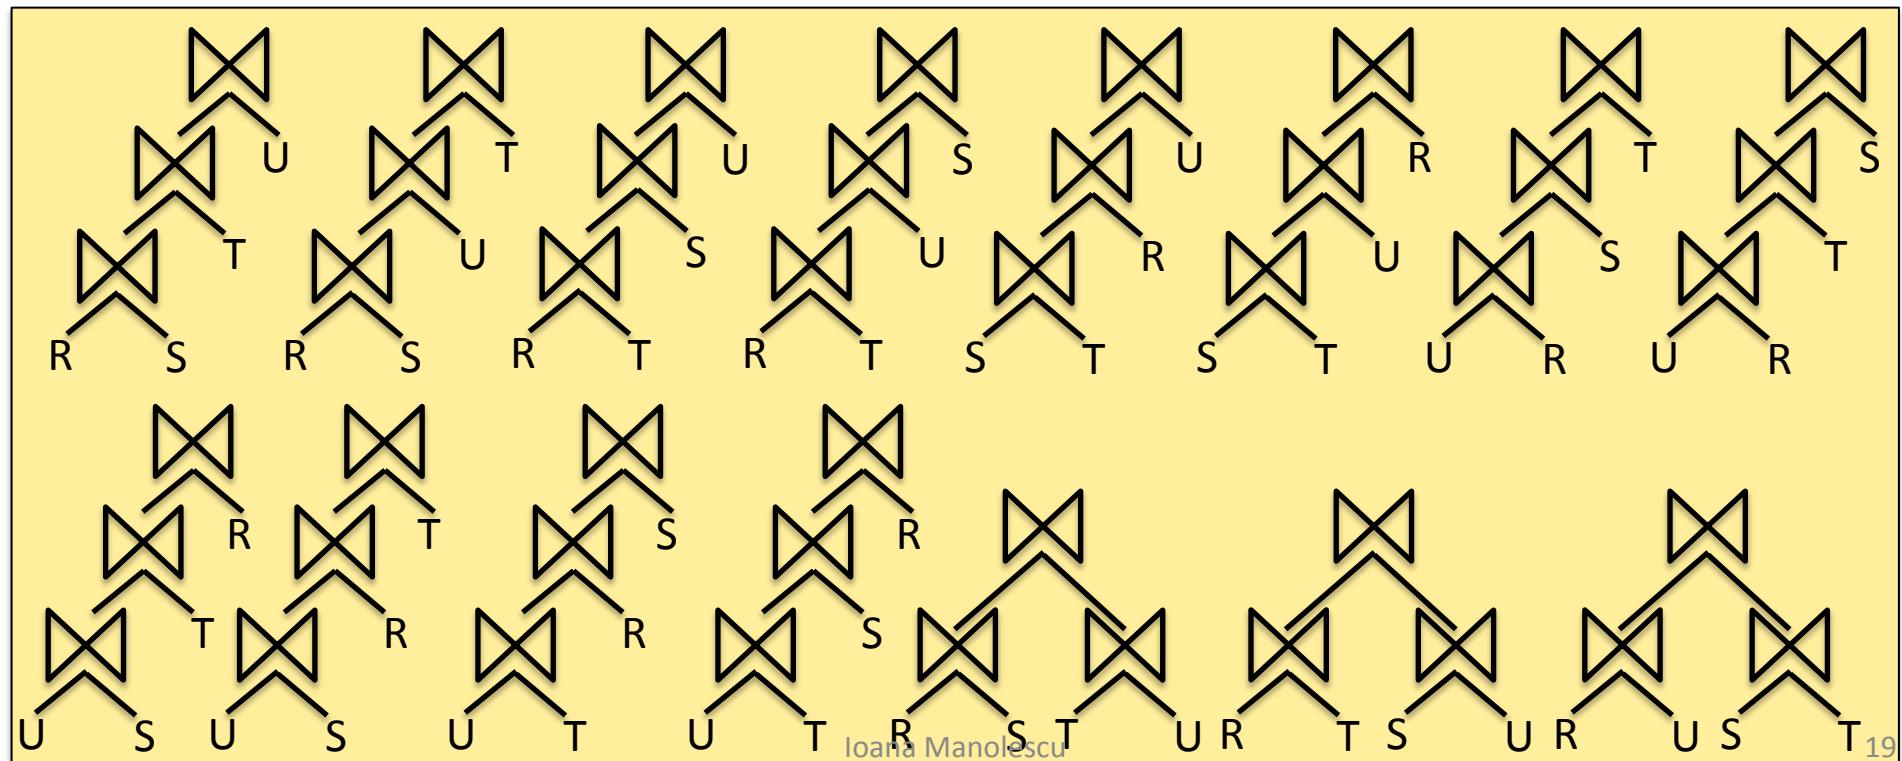
Join ordering is the main problem in logical query optimization

$$\text{Plans}(n+1) = (n+1) * \text{Plans}(n) + \frac{1}{2} * \sum_{i=1}^{(n/2)} \text{Plans}(i) * \text{Plans}(n+1-i)$$

High (exponential) complexity → many heuristics

- Exploring only left-linear plans etc.

N=4:



Logical query optimization needs statistics

Exact statistics (on base data):

- 1.000.000 cars, 1.000.000 drivers, 1.000 accidents

Approximate / estimated statistics (on intermediary results)

- "1.75 cars involved in every accident"

Statistics are gathered

- When **loading** the data: take advantage of the scan
- **Periodically** or upon **request** (e.g. analyze in the Postgres RDBMS)
- At **runtime**: modern systems may do this to change the data layout

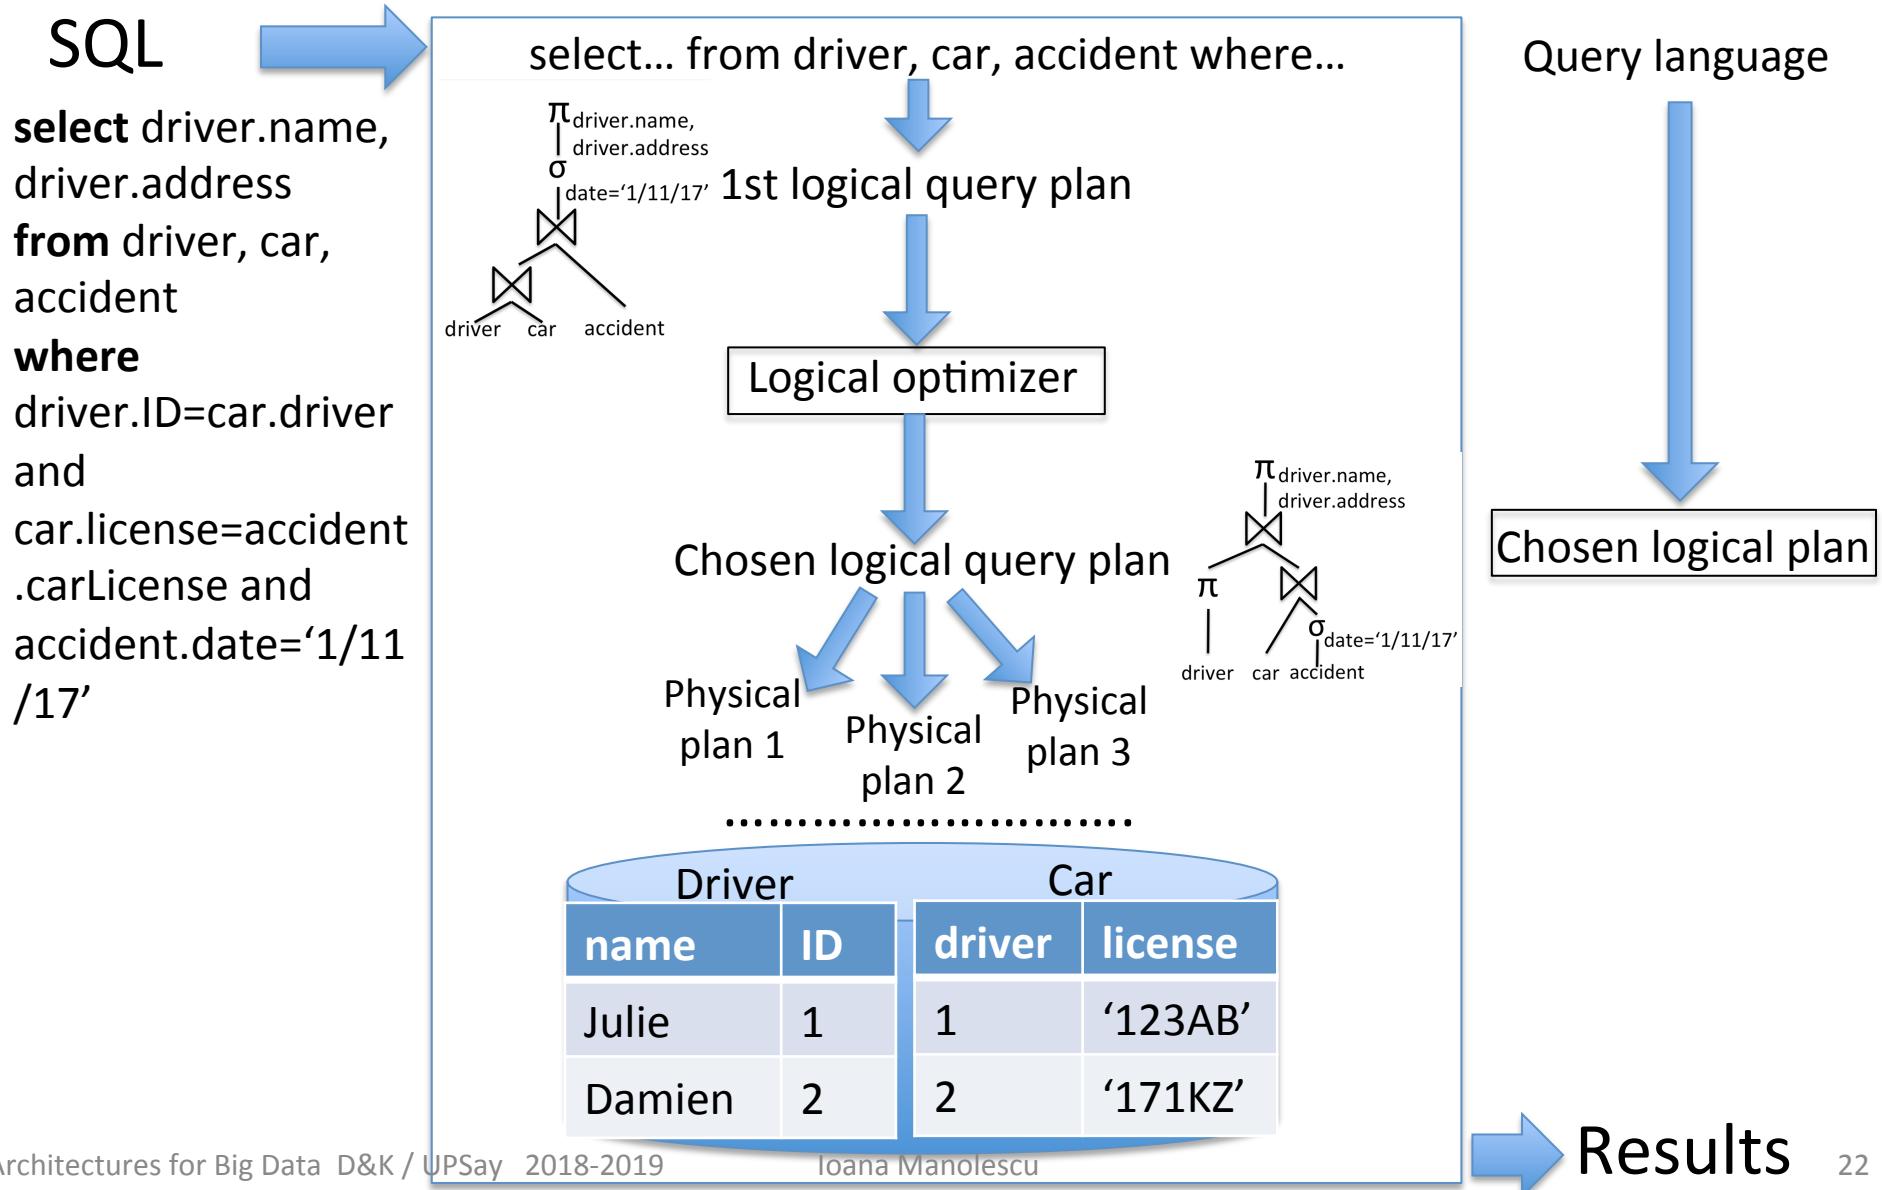
Statistics on the **base data** vs. on **results of operations not evaluated** (yet):

- « On average 2 cars per accident »
- For each column R.a, store:
 $|R|$, $|R.a|$ (number of distinct values), $\min\{R.a\}$, $\max\{R.a\}$
- Assume **uniform distribution** in R.a
- Assume **independent distribution**
 - of values in R.a vs values in R.b; of values in R.a vs values in S.c
- + simple probability computations

More on statistics

- For each column R.a, store:
 $|R|$, $|R.a|$ (number of distinct values), $\min\{R.a\}$, $\max\{R.a\}$
- Assume **uniform distribution** in R.a
- Assume **independent distribution**
 - of values in R.a vs values in R.b; of values in R.a vs values in S.c
- The **uniform distribution** assumption is **frequently wrong**
 - Real-world distribution are skewed (popular/frequent values)
- The **independent distribution** assumption is **sometimes wrong**
 - « Total » counter-example: *functional dependency*
 - Partial but strong enough to ruin optimizer decisions: *correlation*
- Actual optimizers use more sophisticated statistic informations
 - **Histograms**: equi-width, equi-depth
 - Trade-offs: size vs. maintenance cost vs. control over estimation error

Database internal: query optimizer



Physical query plans

Made up of **physical operators** =

algorithms for implementing logical operators

Example: equi-join (R.a=S.b)

Nested loops join:

```
foreach t1 in R{  
    foreach t2 in S {  
        if t1.a = t2.b then output (t1 || t2)  
    }  
}
```

Merge join: // requires sorted inputs

```
repeat{  
    while (!aligned) { advance R or S };  
    while (aligned) { copy R into topR, S into topS };  
    output topR x topS;  
} until (endOf(R) or endOf(S));
```

Hash join: // builds a hash table in memory

```
While (!endOf(R)) { t ← R.next; put(hash(t.a), t); }  
While (!endOf(S)) { t ← S.next;  
    matchingR = get(hash(S.b));  
    output(matchingR x t);  
}
```

Physical query plans

Made up of **physical operators** =

algorithms for implementing logical operators

Example: equi-join ($R.a=S.b$)

Nested loops join:

```
foreach t1 in R{  
    foreach t2 in S {  
        if t1.a = t2.b then output (t1 || t2)  
    }  
}
```

$O(|R| \times |S|)$

Merge join: // requires sorted inputs

```
repeat{  
    while (!aligned) { advance R or S };  
    while (aligned) { copy R into topR, S into topS };  
    output topR x topS;  
} until (endOf(R) or endOf(S));
```

$O(|R| + |S|)$

Hash join: // builds a hash table in memory

```
While (!endOf(R)) { t  $\leftarrow$  R.next; put(hash(t.a), t); }  
While (!endOf(S)) { t  $\leftarrow$  S.next;  
    matchingR = get(hash(S.b));  
    output(matchingR x t);  
}
```

$O(|R| + |S|)$

Also:

Block nested loops join

Index nested loops join

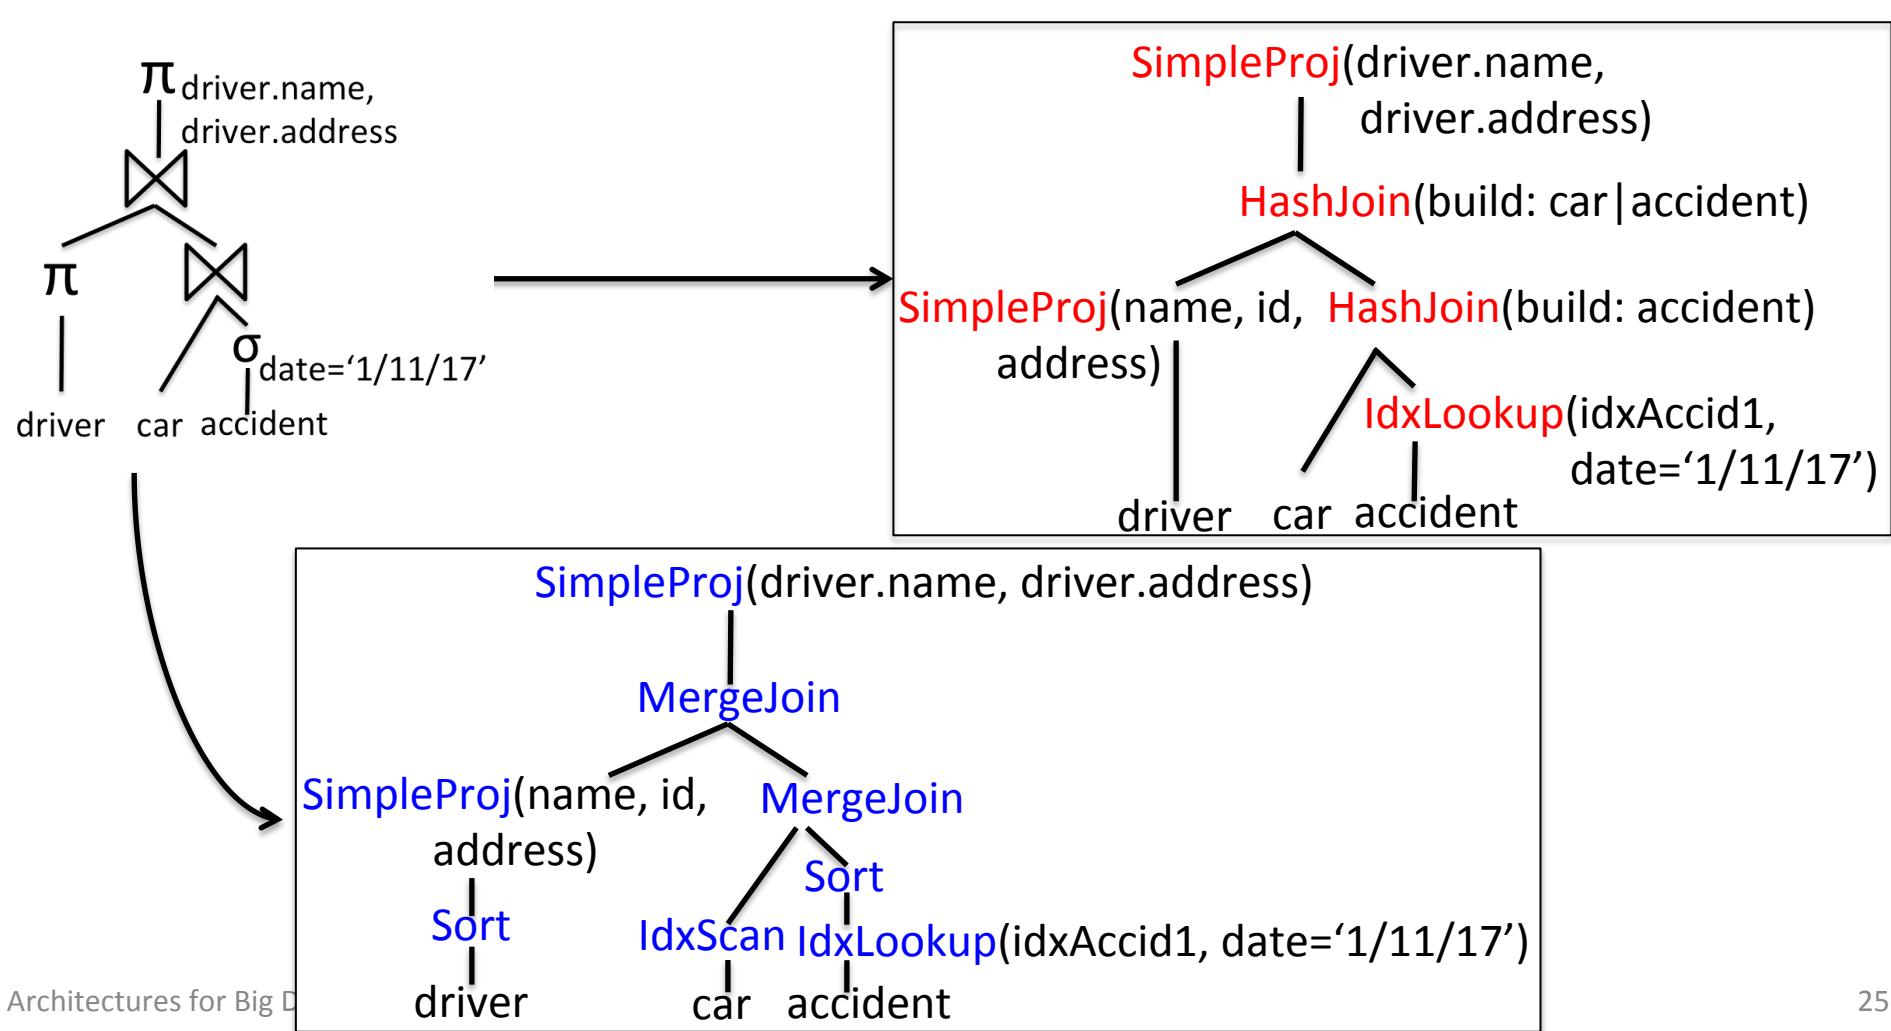
Hybrid hash join

Hash groups / teams

...

Physical optimization

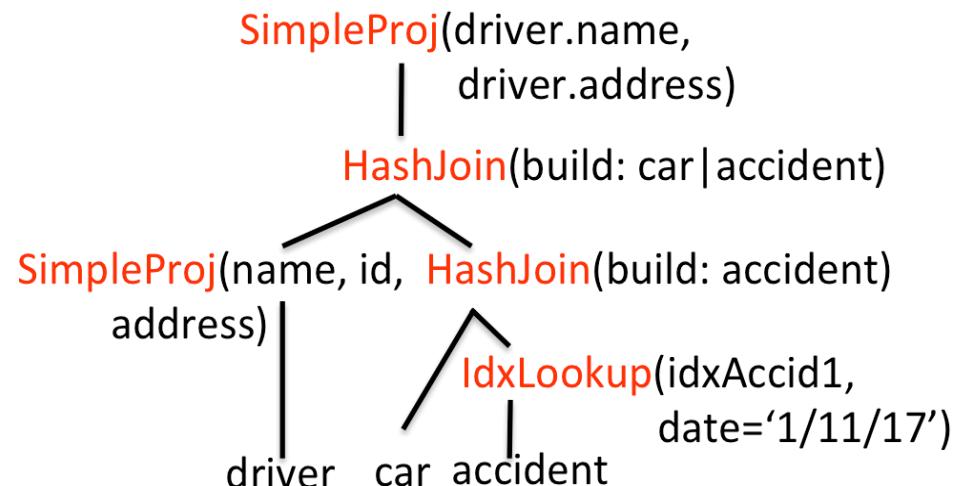
Possible physical plans produced by physical optimization for our sample logical plan:



Physical plan performance

Metrics characterizing a physical plan

- **Response time:** between the time the query starts running to the we know it's end of results
- **Work (resource consumption)**
 - How many **I/O** calls (blocks read)
 - Scan, IdxScan, IdxAccess; Sort; HybridHash (or spilling HashJoin)
 - How much **CPU**
 - All operators
 - Distributed plans: **network** traffic
- **Total work:** work made by all operators



Query optimizers in action

Most database management systems have an « explain » functionality → physical plans. Below sample Postgres output:

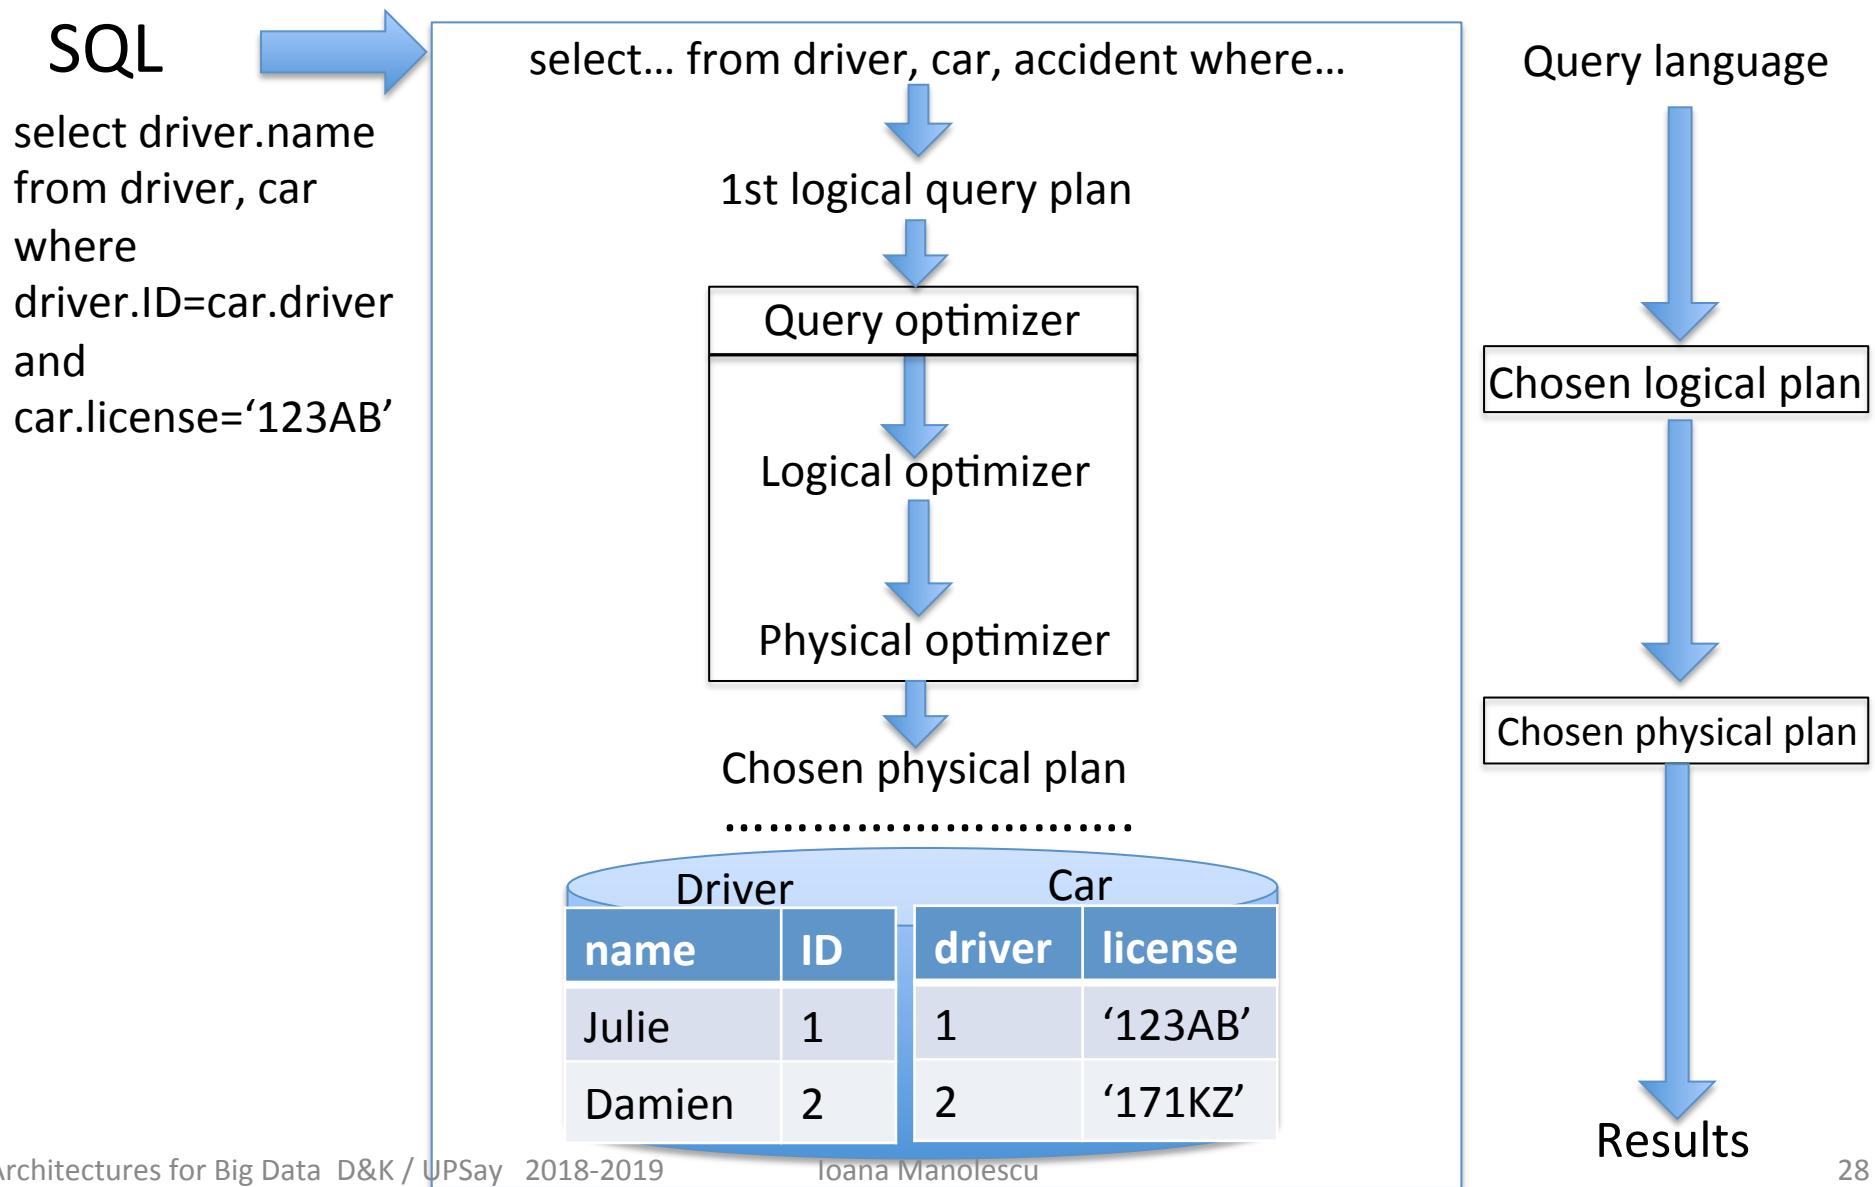
```
EXPLAIN SELECT * FROM tenk1;  
          QUERY PLAN
```

```
Seq Scan on tenk1 (cost=0.00..458.00 rows=10000 width=244)
```

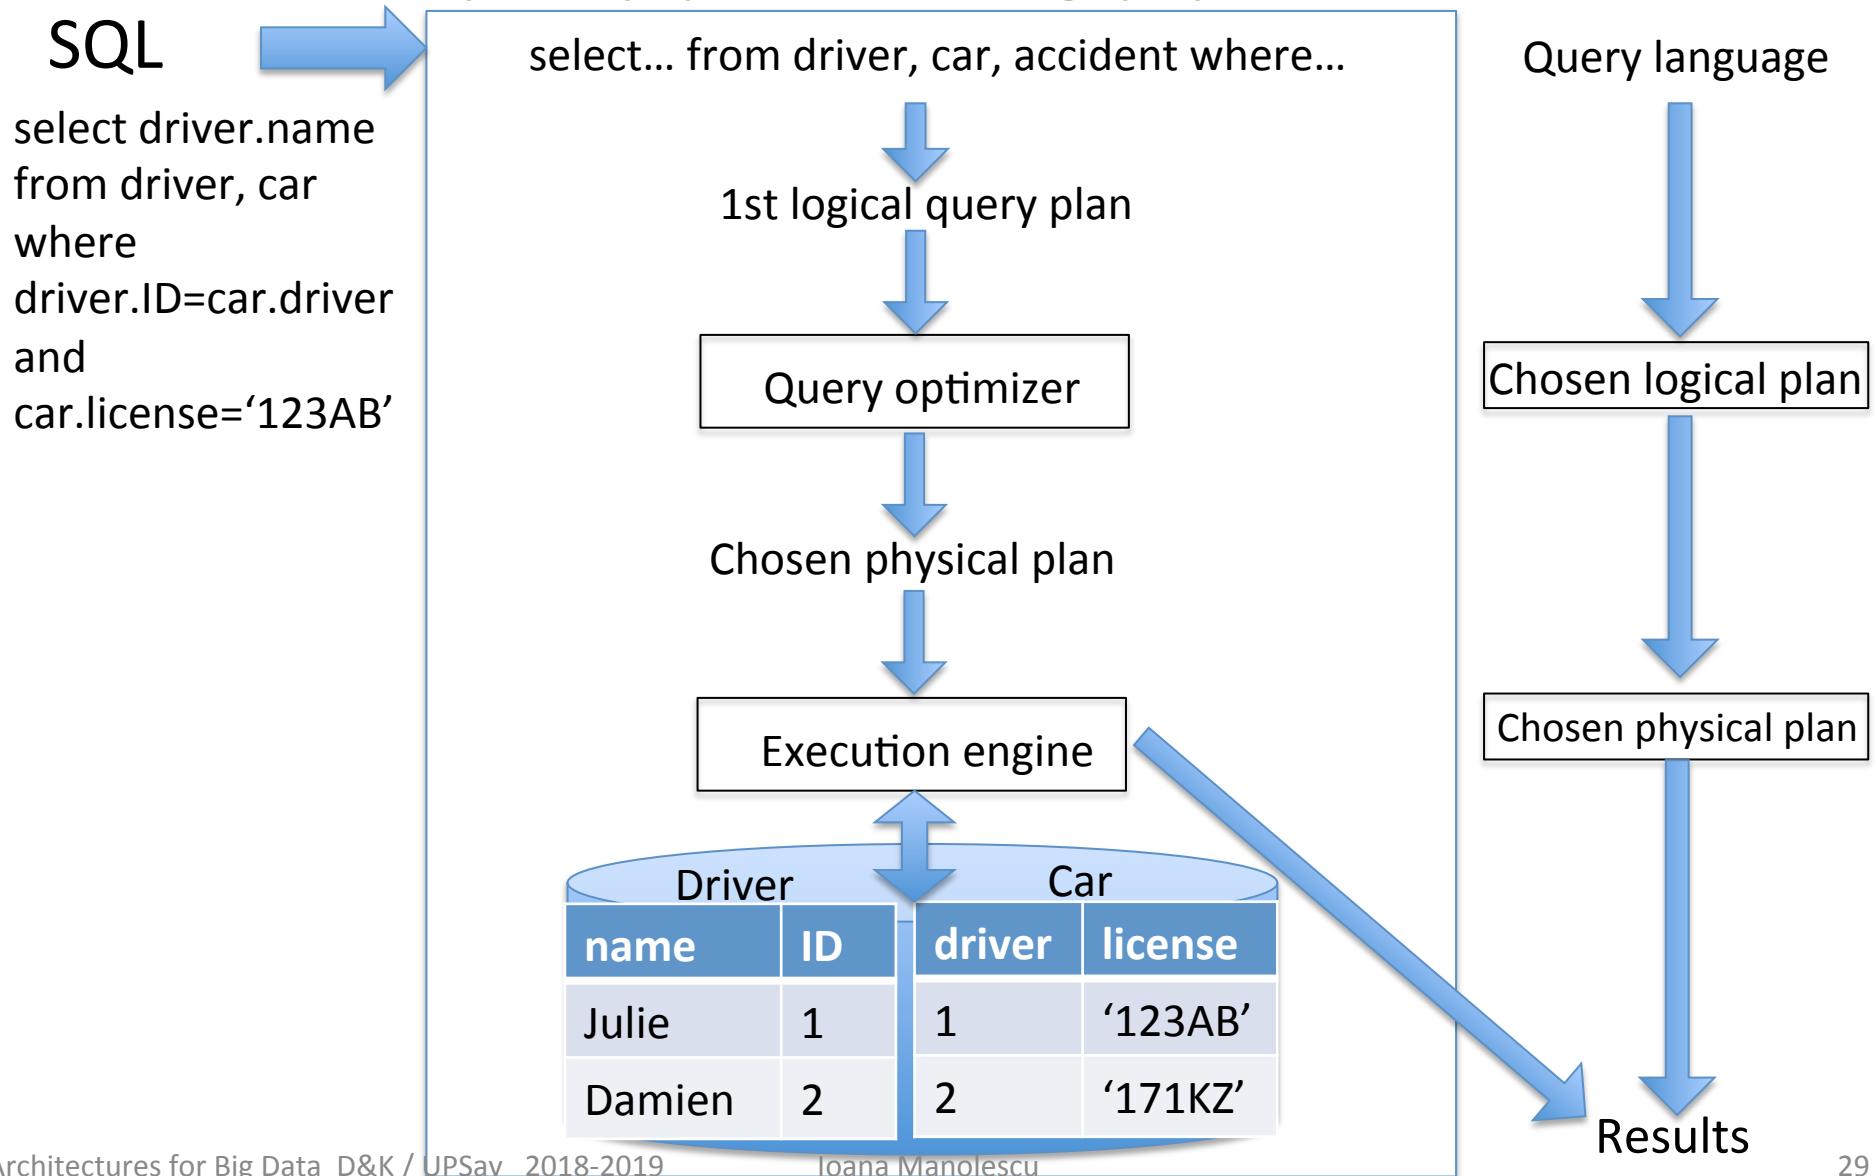
```
EXPLAIN SELECT * FROM tenk1 t1, tenk2 t2  
WHERE t1.unique1 < 100 AND t1.unique2 = t2.unique2;  
          QUERY PLAN
```

```
Hash Join (cost=232.61..741.67 rows=106 width=488)  
  Hash Cond: ("outer".unique2 = "inner".unique2)  
  -> Seq Scan on tenk2 t2 (cost=0.00..458.00 rows=10000 width=244)  
  -> Hash (cost=232.35..232.35 rows=106 width=244)  
      -> Bitmap Heap Scan on tenk1 t1 (cost=2.37..232.35 rows=106 width=244)  
          Recheck Cond: (unique1 < 100)  
          -> Bitmap Index Scan on tenk1_unique1 (cost=0.00..2.37 rows=106 width=0)  
              Index Cond: (unique1 < 100)
```

Database internal: physical plan



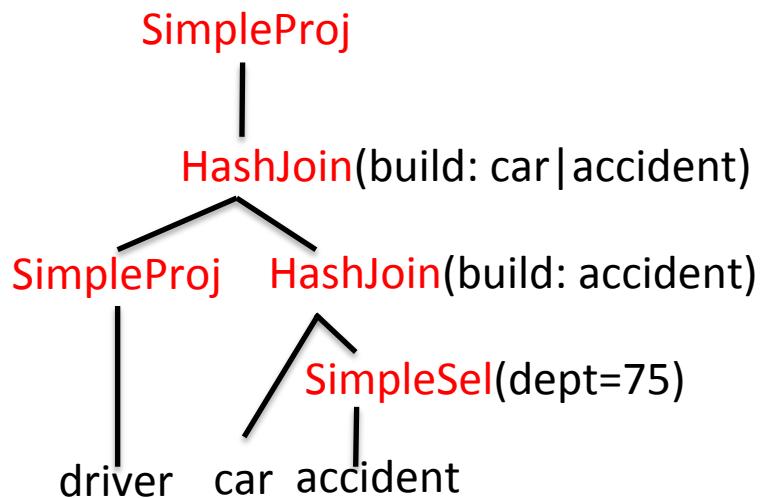
Database internals: query processing pipeline



Advanced query optimization techniques: Dynamic Query Optimization

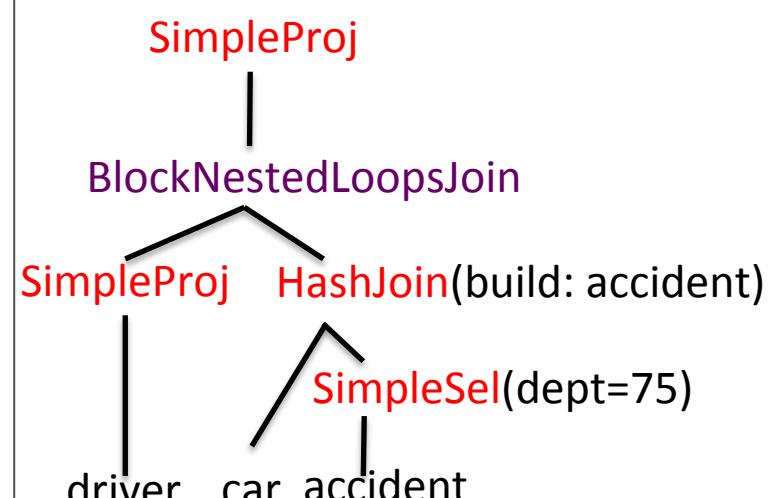
- Sizes (cardinalities) of intermediary results are estimated, which may lead to estimation errors
- A cardinality estimation error may lead to choosing a logical plan and a set of physical operators that perform significantly different from expectation (especially for the worse)

Initially chosen plan:



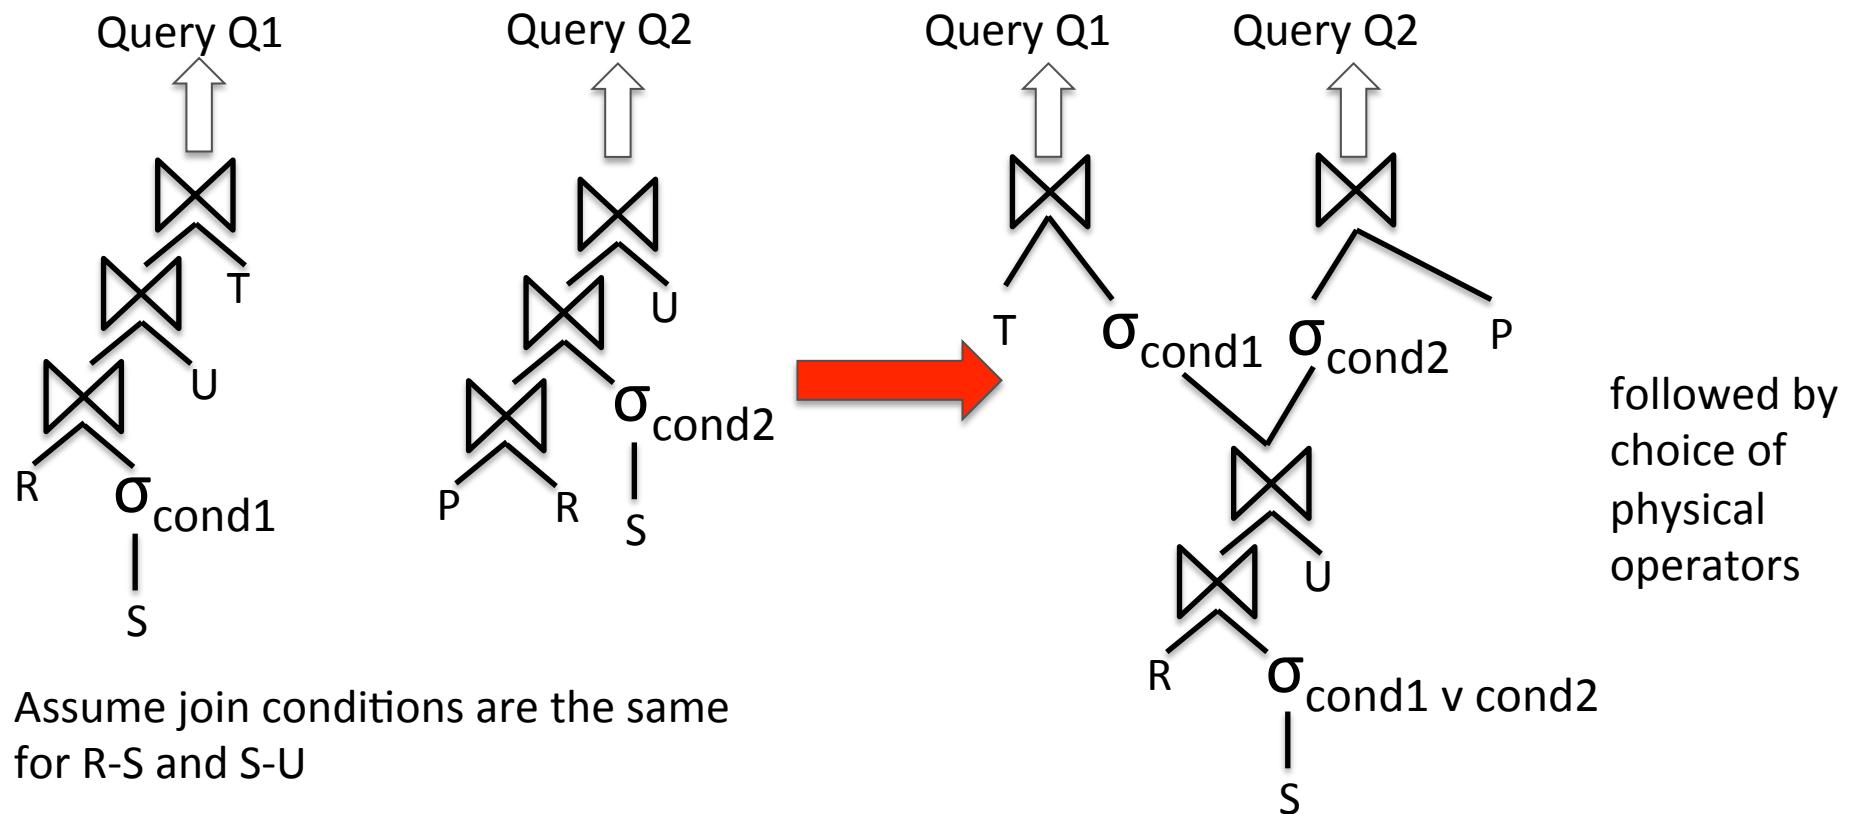
At execution time, we see that the lower HashJoin output is larger than expected: memory insufficient to build

Modified plan:



Advanced query optimization techniques: Multi-Query Optimization

Multiple queries sharing sub-expressions can be optimized together into a single plan with **shared subexpressions**



What's in a database?

SQL update

```
insert into driver  
values ('Thomas',  
3);  
update car set  
driver=3 where  
license='123AB';
```

Database

Driver		Accident		Car	
name	ID	driver	license		
Julie	1	1	'123AB'		
Damien	2	2	'171KZ'		

Database

Driver		Accident		Car	
name	ID	driver	license		
Julie	1	3	'123AB'		
Damien	2	2	'171KZ'		
Thomas	3				

Database updates

- A set of operations atomically executed (either all, or none) is called a **transaction**
- There may be some **dependencies** between the operations of a transaction
 - First read the bank account balance
 - Then write that value reduced by 50€
- A total order over the operations of several concurrent transaction is called a **scheduling**
- The DB component that receives all incoming transactions and decides what operation will be executed when (i.e., global order over the operations of all transactions) is called a **scheduler**

Database updates

- The scheduler is in charge of ordering all operations so that they will appear executed one after the other (serially)

```
T1: BEGIN A=A+100, B=B-100 END
```

```
T2: BEGIN A=1.06*A, B=1.06*B END
```

```
T1: A=A+100, B=B-100,
```

```
T2: A=1.06*A, B=1.06*B
```

```
T1: A=A+100, B=B-100
```

```
T2: A=1.06*A, B=1.06*B
```

ARCHITECTURES FOR BIG DATA:

WHAT NEEDS TO CHANGE?

What is the impact of Big Data properties on database architectures?

- **Volume** requires **distribution**
 - Of the data storage; of query evalution
 - Distribution makes **ACID difficult** (CAP theorem) ✓
 - Complicates concurrency control
 - Replication, eventual consistency
 - Distribution requires efficient, easy-to-use **parallelism**
 - Distribution raises issues of **control** which can lead to single point of failure → **decentralization**
- **Velocity** requires efficient algorithms
 - Optimize for throughput (rather than response time)
 - Stream processing, in-memory architectures
 - Process-then-store (or process-then-discard)

What is the impact of Big Data properties on database architectures?

- **Variety** requires support for
 - **flexible data models**: key-values, JSON, graphs...
 - **different schemas**, and translation mechanisms between the schemas
 - Data integration
 - **several data models** being used together
 - Mediators, Data lakes
- **Veracity** requires support for reconciliation, data cleaning etc.
 - Similar to single-database setting, but adding source, source confidence, and provenance information

Roadmap for the rest of the course

1. Analysis of large-scale (in particular, distributed) Big Data platforms
 - Focus on: **distribution** of data and query processing, concurrency control
2. A selection of NoSQL platforms
 - Choice of most popular ones in their class
 - To illustrate **variety** of data models, and some distribution choices